


Object-Oriented Programming

Programming Project #2+3

Data Center

- A data center consists of multiple servers
- The servers are connected by switches in a local area network




Switches

- Each switch has multiple ports
- Receive and forward the packets from a port to another port




Traditional Routing Path

- Switches use OSPF (i.e., shortest path)
- Construct a shortest path tree rooted at each destination



OSPF Routing Information


- Given: a graph with links and destinations
- Output: shortest paths towards all destinations
- Then, store the information in each node's table



OSPF Routing Table

- Key: each **destination**
- Value: the **next node** (i.e., the output port)
- **Node 1's table** (it uses OSPF)

Destination	Next Node
0	2




Programming Project #2:

Implement OSPF in Traditional Networks

- Input:
 - # nodes, # destinations, # links, and #pairs
 - Destinations (ID)
 - Links between nodes
 - Traffic matrix (flow size for each pair)
- Procedure:
 - Compute shortest paths to each destination
using distributed BFS with a counter
(choose the node with a smaller ID if there is a tie)
 - Forward packets with the help of routing tables
- Output:
 - Every node's routing table
 - Packet exchange information will be logged automatically


Note – Rule for Selecting the Next Hop

- Select the node with **a smaller counter** (i.e., closer to the destination)
- Select the node with **a smaller ID** as the next hop **if there is a tie** (i.e., multiple candidates)



Note – Rule for Relaying TRA_ctrl_packet

- Relay packets with a counter smaller than all my currently received counters
- Relay packets with a counter equal to my parent's counter but with a preID smaller than my current parent



Note – Create TRA Nodes and Links

- Create traditional switches (i.e., `class TRA_switch`)
- Each switch has an unsigned int `ID`
 - `node::node_generator::generate("TRA_switch", id);`
- Every node only `knows its neighbors`
- `Add the neighbors` for each switch
 - `node::id_to_node(0)->add_phy_neighbor(1);`
 - `node::id_to_node(1)->add_phy_neighbor(0);`
 - We use `simple_link` with a fixed latency (i.e., 10)
- Write a `map<unsigned int, unsigned int>` to store `each entry` in each switch's table (i.e., each entry in the table has `<a destination ID, a next node ID>`) in class `TRA_switch`
 - Copy and modify partial code in HW1

Note – Generate Data and Ctrl Packets

- **Generate data packets**

- void **data_packet_event**(unsigned int src, unsigned int dst, unsigned int t = 0, string msg="default")
- A TRA_data_packet will be **generated for a source (src)** and **sent to a destination (dst) at time t**
- The source (src) will receive the TRA_data_packet first (since it's src)

- **Generate ctrl packets**


- void **TRA_ctrl_packet_event**(unsigned int src, unsigned int t = event::getCurTime(), string msg = "default")
- The function is used to initialize the distributed BFS; that is, a TRA_ctrl_packet will be **generated for a source (src) with a counter 0**
- You have to implement **recv_handler()** in TRA_switch to forward the ctrl packet to the neighboring node again; that is, **every node receiving the packet should increase the counter and broadcast the packet to its neighboring nodes** to update the rule in every node's table to build the path from every node to the source
- The source (src) will receive the TRA_ctrl_packet first (since it's src)

Note – Receive and Send Packets (1/2)

- Define the rules to **handle the received packet** in class TRA_switch's member function recv_handler
 - void TRA_switch::recv_handler (packet *p)
 - **Don't use** node::id_to_node(id) in recv_handler
- Get the **current switch's ID and its neighbor**
 - Use getNodeID() in recv_handler
 - Use getPhyNeighbors().find(n_id) to check whether the node with n_id is a neighbor
 - Use const map<unsigned int,bool> &nblast = getPhyNeighbors() and for (map<unsigned int,bool>::const_iterator it = nblast.begin(); it != nblast.end(); it++) to get all neighbors
- Use **send_handler(packet *p)** to send the packet
- Check the packet type
 - if (p->type() == "TRA_data_packet")
 - if (p->type() == "TRA_ctrl_packet")

Note – Receive and Send Packets (2/2)

- Decode: **Cast the packet, payload**, to the right type
 - `TRA_data_packet *p2 = dynamic_cast<TRA_data_packet*> (p)`
 - `TRA_ctrl_packet *p3 = dynamic_cast<TRA_ctrl_packet*> (p)`
 - `TRA_ctrl_payload *l3 = dynamic_cast<TRA_ctrl_payload*> (p3->getPayload());`
 - ...
 - Will be explained in the later chapters
- **Before sending** a packet to the next hop
 - Use `setPreID(id)` to change the preID to the current node's ID
 - Use `setNexID(id)` to change the nexID to the next hop node's ID
 - Please check all the columns in the header

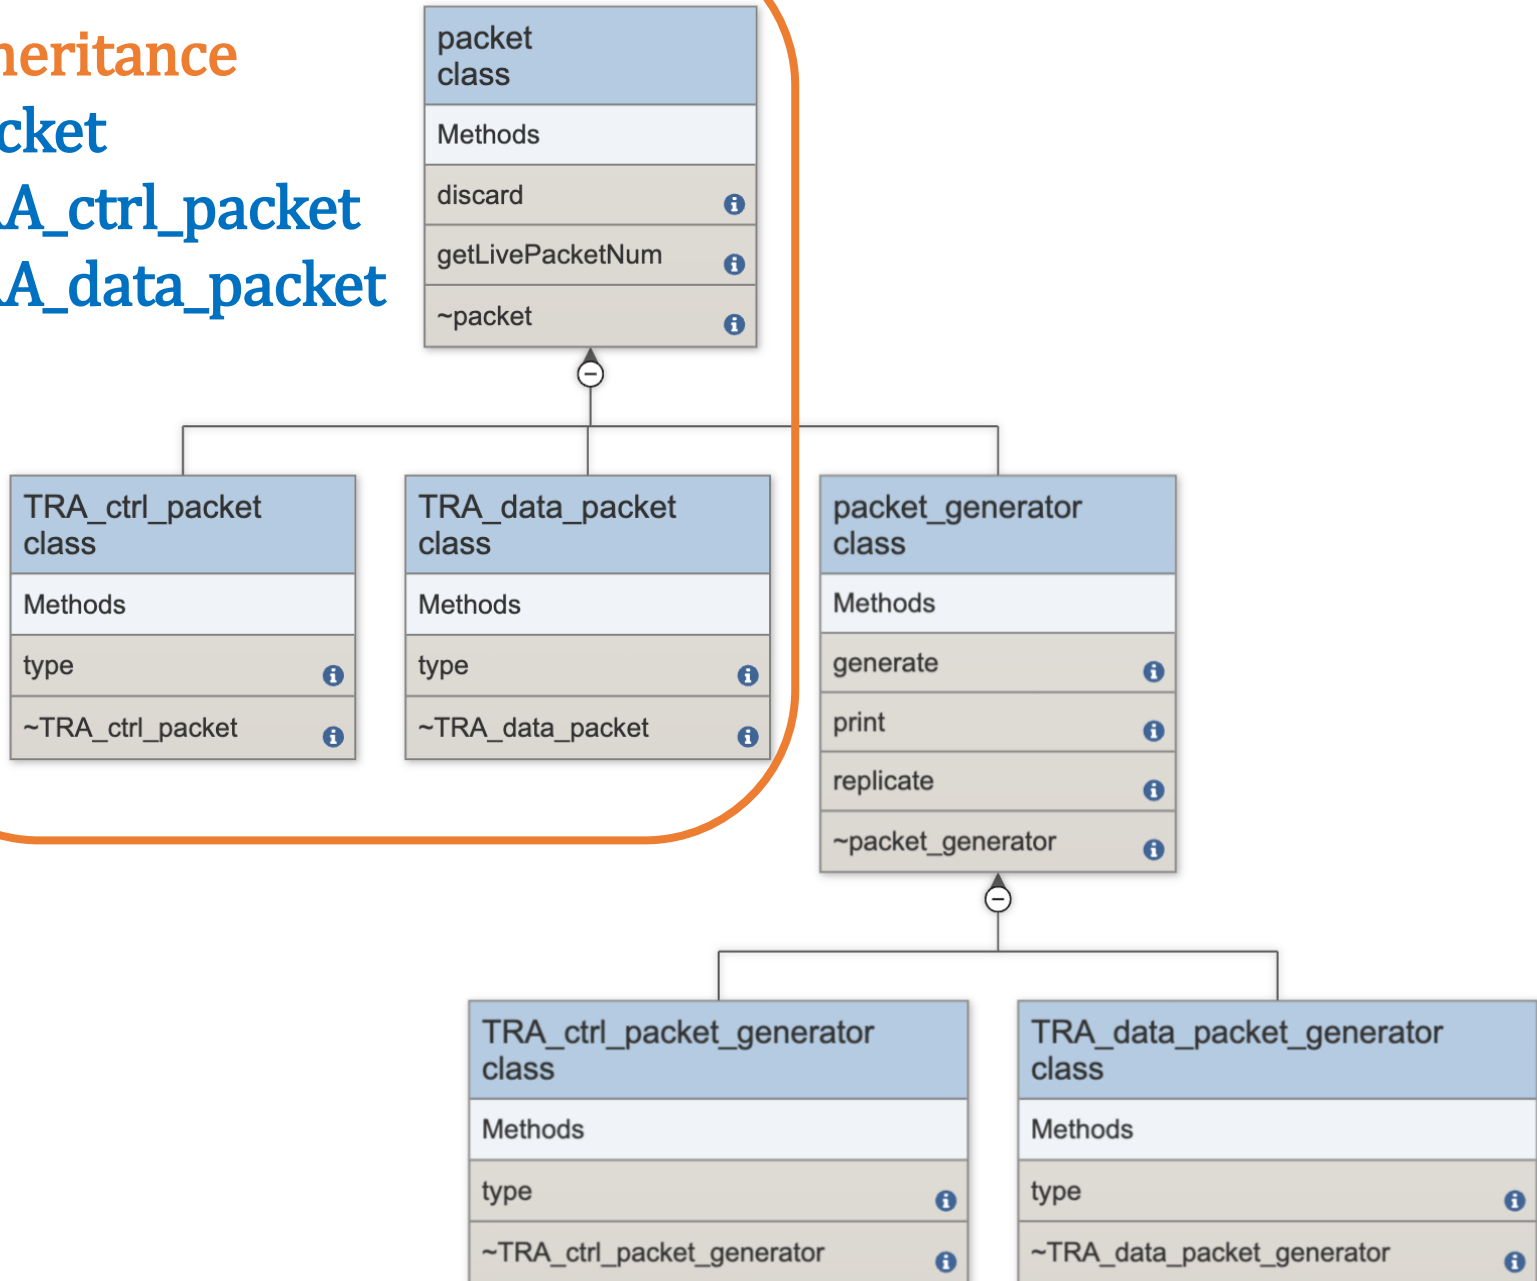


Inheritance

packet

TRA_ctrl_packet

TRA_data_packet

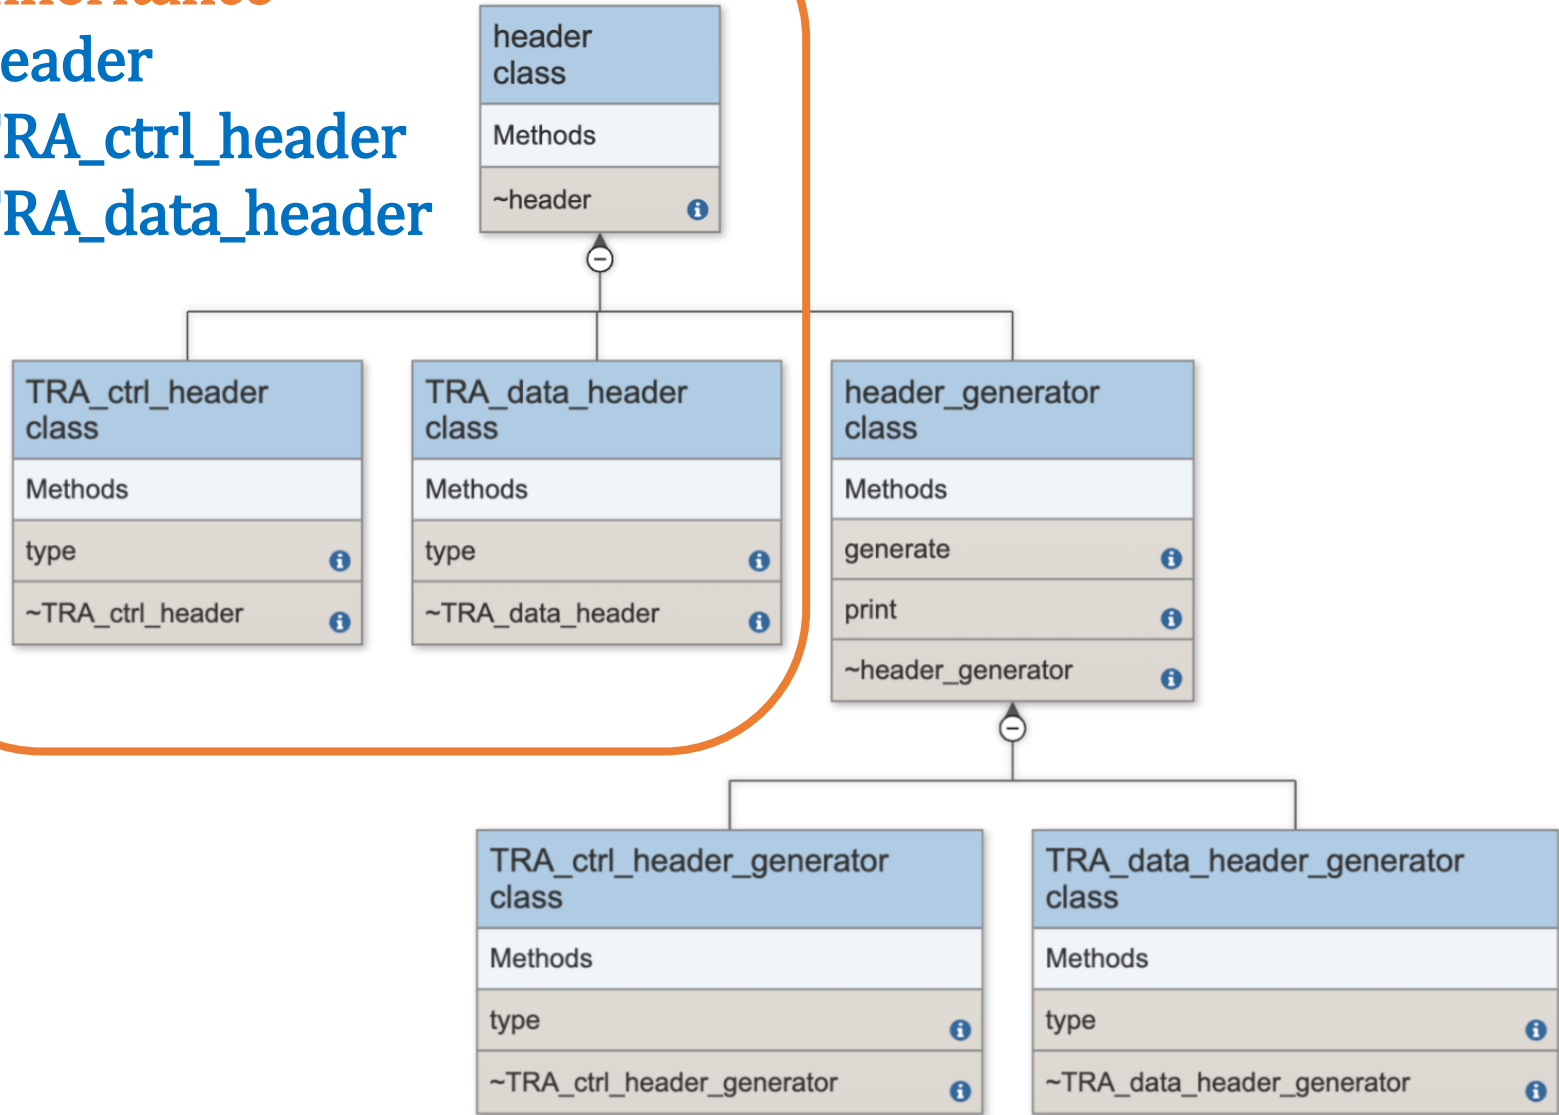


Inheritance

header

TRA_ctrl_header

TRA_data_header

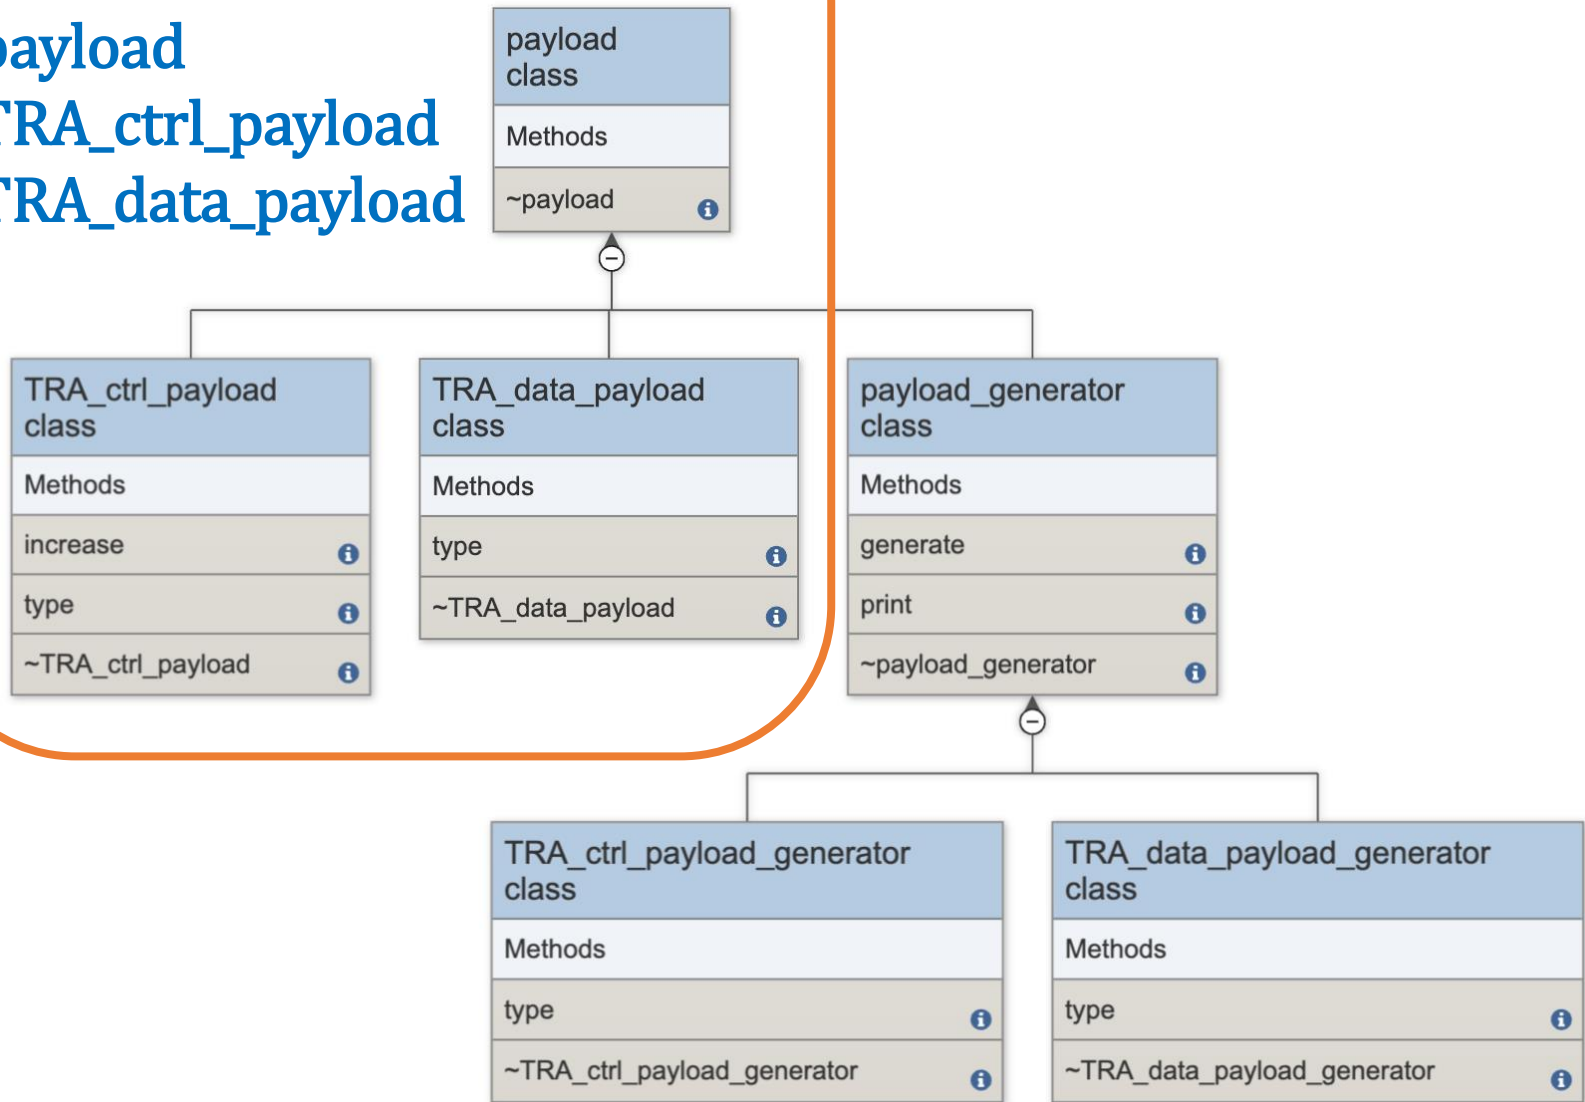


Inheritance

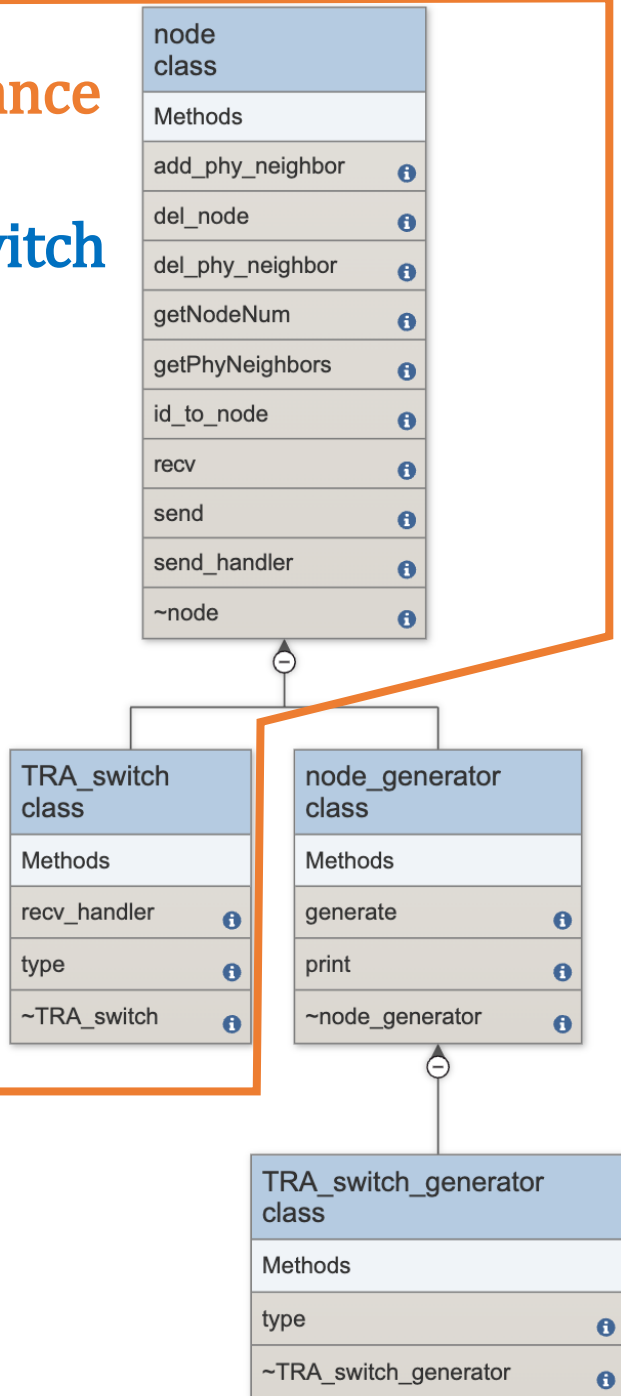
payload

TRA_ctrl_payload

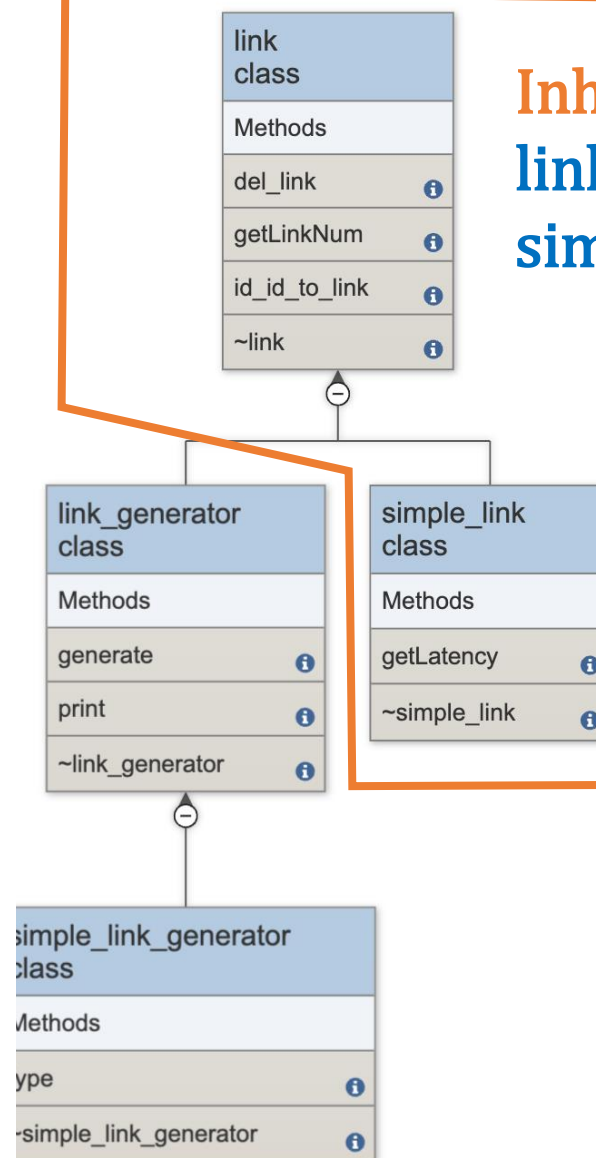
TRA_data_payload



Inheritance node TRA_switch



Inheritance link simple_link



Input Sample:

use cin

15 1 28 3 300

0

0 100

Format:

#Nodes #Dsts #Links #Pairs SimTime

DstID_List

NodeID BroadcastTime


...

LinkID Node1 Node2

...

FlowID Src Dst FlowSize StartTime

...



0 0 5
1 0 10
2 0 12
3 0 14
4 1 2
5 1 3
6 1 4
7 2 3
8 2 10
9 2 11
10 3 4
11 3 6
12 3 7
13 4 8
14 4 9

15 5 6
16 5 7
17 5 9
18 5 10
19 5 14
20 6 7
21 6 10
22 7 9
23 8 9
24 9 13
25 10 12
26 11 12
27 13 14

0 1 0 3 150
1 2 0 4 200
2 3 0 5 120

Output Sample:

use cout

You have to print the routing table for each node after event::start_simulate();

The way to print the routing table is the same as that in HW1

The remaining output will be automatically generated 😊

Note that the output could be different in different computers

Its own ID

Output Sample (continue):

use cout

Format:

The automatic printing (you can't change)


NodeID


DstID NextID

...

e.g.,


0		7
0	0	0 5
1		8
0	2	0 9
2		9
0	10	0 5
3		10
0	2	0 0
4		11
0	9	0 12
5		12
0	0	0 0
6		13
0	5	0 14
		14
		0 0







Traditional Switches vs SDN switches

- OSPF paths are **fixed** → Not flexible
- SDN-enabled switches provide **path diversity**
- However, SDN switches are **expensive**
- We cannot upgrade all switches



Which Switch Should be Upgraded First?

- Given a **limited budget** for upgrading switches
- Each switch has a different **upgrade cost**
- Goal: **minimize the maximum link load**



Programming Project #3:

Partially Upgrading an ISP network to SDN

- Input:
 - # nodes, # destinations, # links, and #pairs
 - Nodes (ID) with its upgrade cost to support SDN
 - Destinations (ID)
 - Links between nodes
 - Traffic matrix (flow size for each pair)
 - Limited budget to bound the total upgrade cost of nodes
- Procedure:
 - Compute a set of nodes that will be upgraded to support SDN
 - Compute shortest paths to each destination
 - Compute next hops and portions for SDN-enabled nodes
- Output:
 - The set of SDN-enabled nodes
 - Each node's routing table
 - Packet exchange information will be logged automatically

The Competition

- The grade is inversely proportional to **the max link load**
- **Basic: 60 (deadline)**
 - **Every node's packet can be sent to the destination with no cycle**
- **Being a coding assistant (superb deadline)**
 - +10
- **Performance ranking (decided after the deadline)**
 - [0%, 30%) (bottom): +0
 - [30%, 50%): + 5
 - [50%, 75%): + 10
 - [75%, 85%): + 15
 - [85%, 90%): + 20
 - [90%, 95%): + 25
 - [95%, 100%] (top): + 30

The Competition Rules

- Note that you **cannot** use brute-force algorithm
- Your solution should be **deterministic** on our server
 - E.g., the random seed & the number of iterations are fixed

Deterministic!



We have a **strict**
TIME LIMIT!

Note – Create SDN Nodes and Links

- Define class **SDN_switch** and Create **SDN_switch**
 - Derived from **class node (Inheritance)**
- Each switch has an unsigned int **ID**
 - `node::node_generator::generate("SDN_switch",id);`
- Every node only **knows its neighbors**
- **Add the neighbors** for each switch
 - `node::id_to_node(0)->add_phy_neighbor(1);`
 - `node::id_to_node(1)->add_phy_neighbor(0);`
 - We use `simple_link` with a fixed latency (i.e., 10)
- Write a **`map<unsigned int, vector<pair<unsigned int, double>>>`** to store **each entry** in each switch's table (i.e., each entry in the table has destination ID, <next nodes' IDs, portions>) in class **SDN_switch**
 - Copy and modify the routing table code in HW1

Note – Receive and Send Packets (1/2)

- Define the rules to **handle the received packet** in class SDN_switch's member function `recv_handler`
 - `void SDN_switch::recv_handler (packet *p)`
 - **Don't use** `node::id_to_node(id)` in `recv_handler`
- Get the **current switch's ID and its neighbor**
 - Use `getNodeID()` in `recv_handler`
 - Use `getPhyNeighbors().find(n_id)` to check whether the node with `n_id` is a neighbor
 - Use `const map<unsigned int,bool> &nblast = getPhyNeighbors()` and `for (map<unsigned int,bool>::const_iterator it = nblast.begin(); it != nblast.end(); it ++)` to get all neighbors
- Use **`send_handler(packet *p)`** to send the packet
- Check the packet type
 - `if (p->type() == "TRA_data_packet")`
 - `if (p->type() == "SDN_ctrl_packet")`

Note – Define and Create SDN Controller

- Define new **class SDN_controller**
 - Derived from **class node (Inheritance)**
- After creating the switches, create an **SDN controller**
 - con_id is the controller ID (after all switches' ID)
 - node_generator::generate("SDN_controller", con_id);
- **Connect each SDN switches** to the controller
 - node::id_to_node(switch_id)->add_phy_neighbor(con_id);
 - node::id_to_node(con_id)->add_phy_neighbor(switch_id);

Note

- **Generate ctrl packets**

- `void SDN_ctrl_packet_event(unsigned int con_id, unsigned int id, unsigned int mat, unsigned int act, double per, unsigned int t = event::getCurTime(), string msg = "default")`
- A packet will be **generated for the controller** and sent to the node (id) **at time t (optional)** to update a specific rule in the node's table
- **mat**: the destination in the node's routing table
act: the next hop in the node's routing table
per: the portion of flow sent to the next node
- The node can use **getMatID(), getActID(), and getPer()** of `SDN_ctrl_payload` to get mat, act, and per from the `SDN_ctrl_packet`
- The controller will receive the `SDN_ctrl_packet` first (since the controller is src)

Inheritance

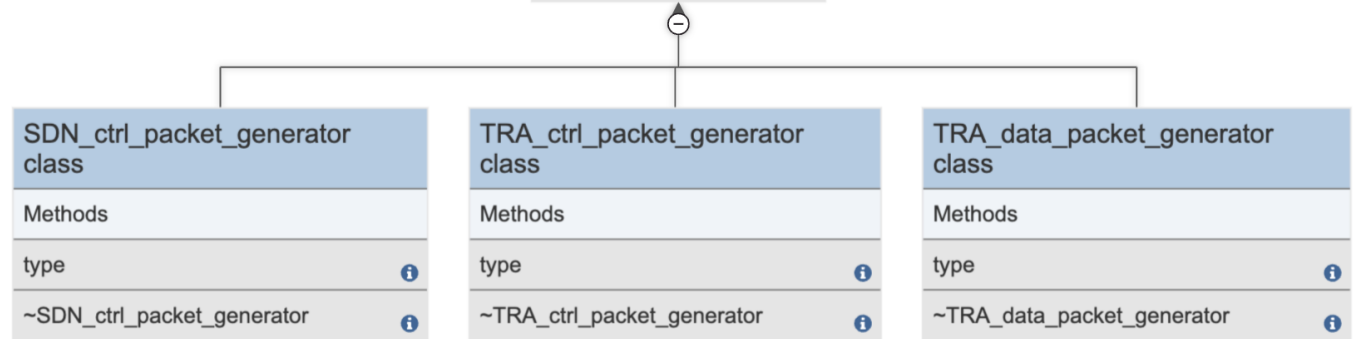
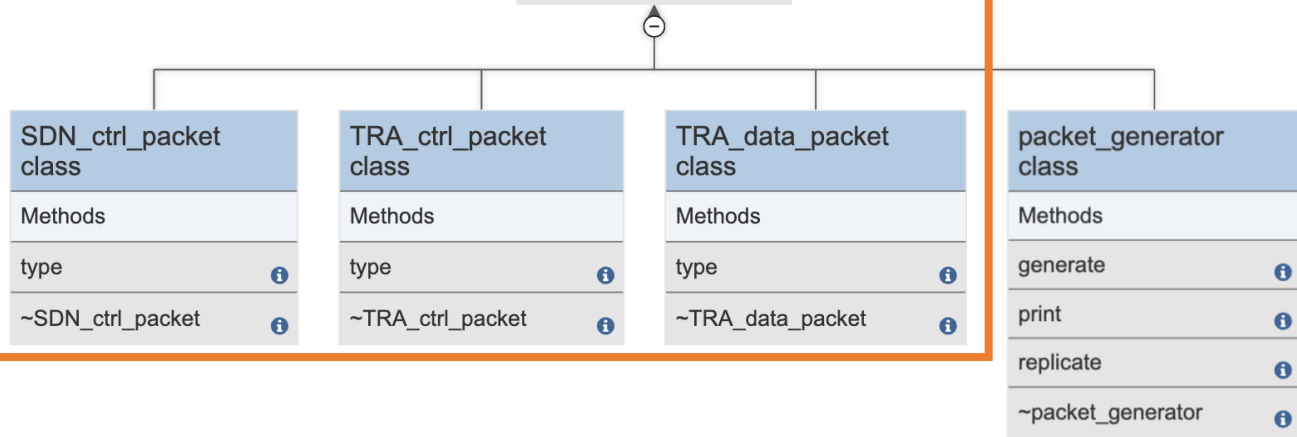
packet

TRA_ctrl_packet

TRA_data_packet

SDN_ctrl_packet

packet class	
Methods	
discard	i
getLivePacketNum	i
~packet	i



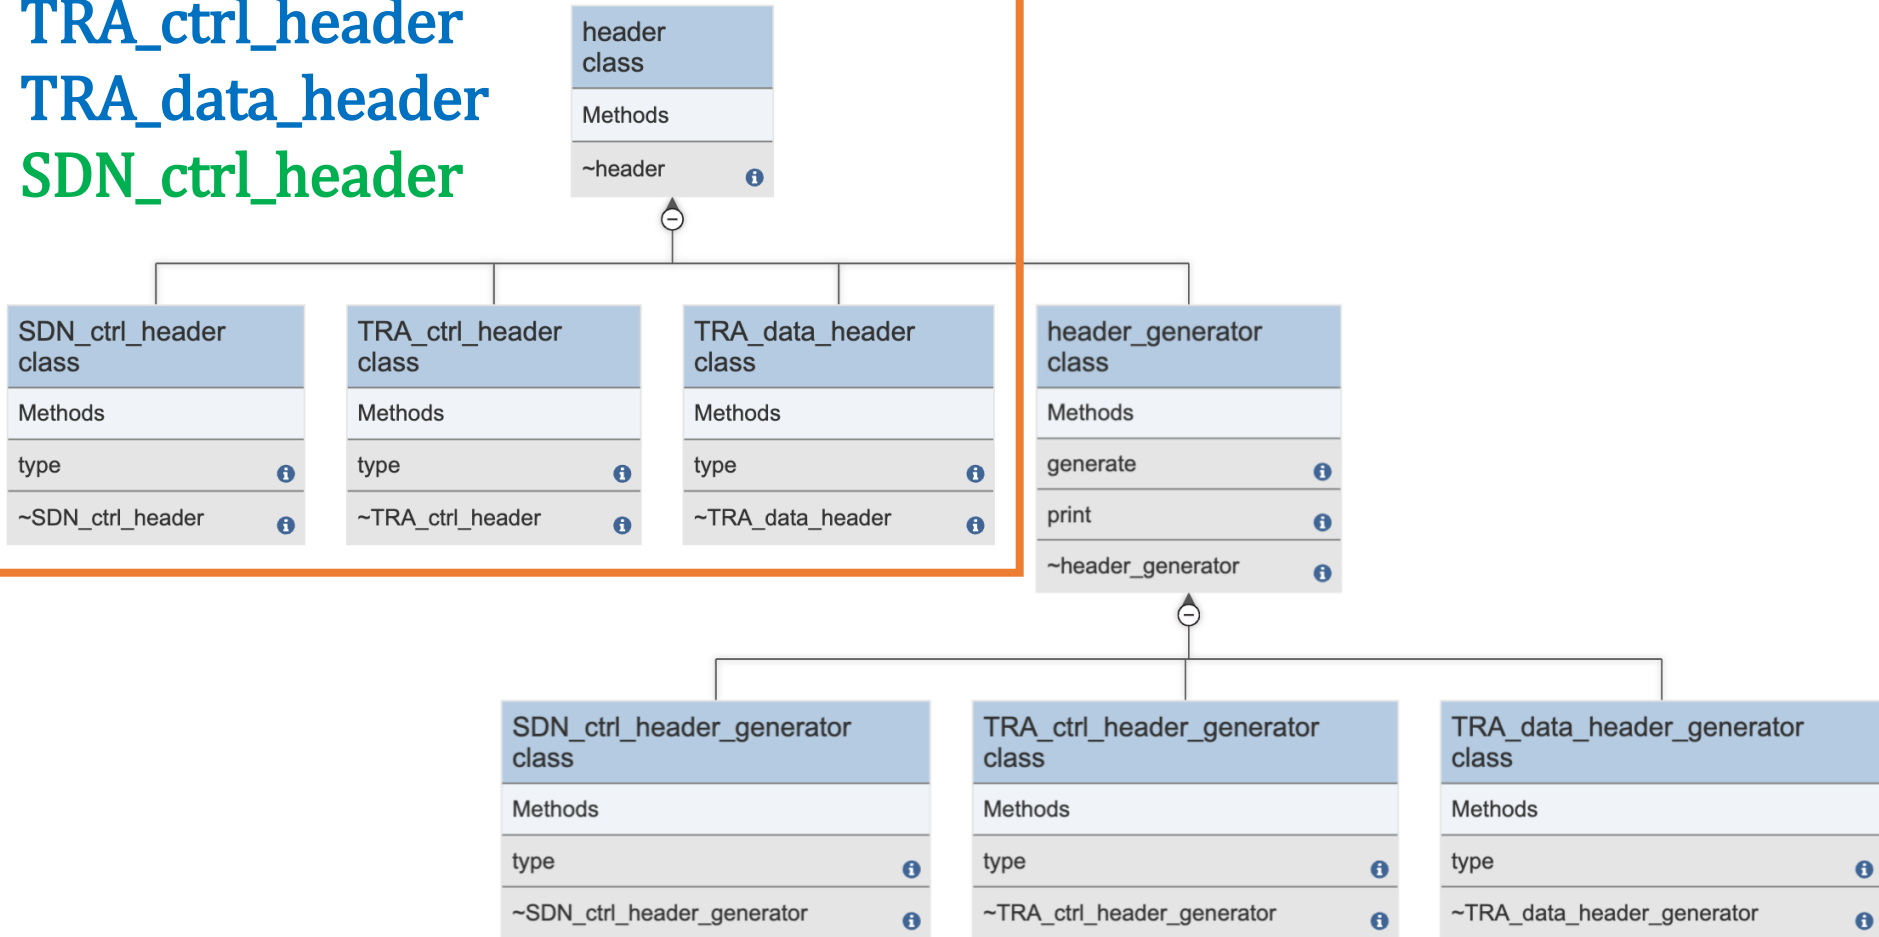
Inheritance

header

TRA_ctrl_header

TRA_data_header

SDN_ctrl_header



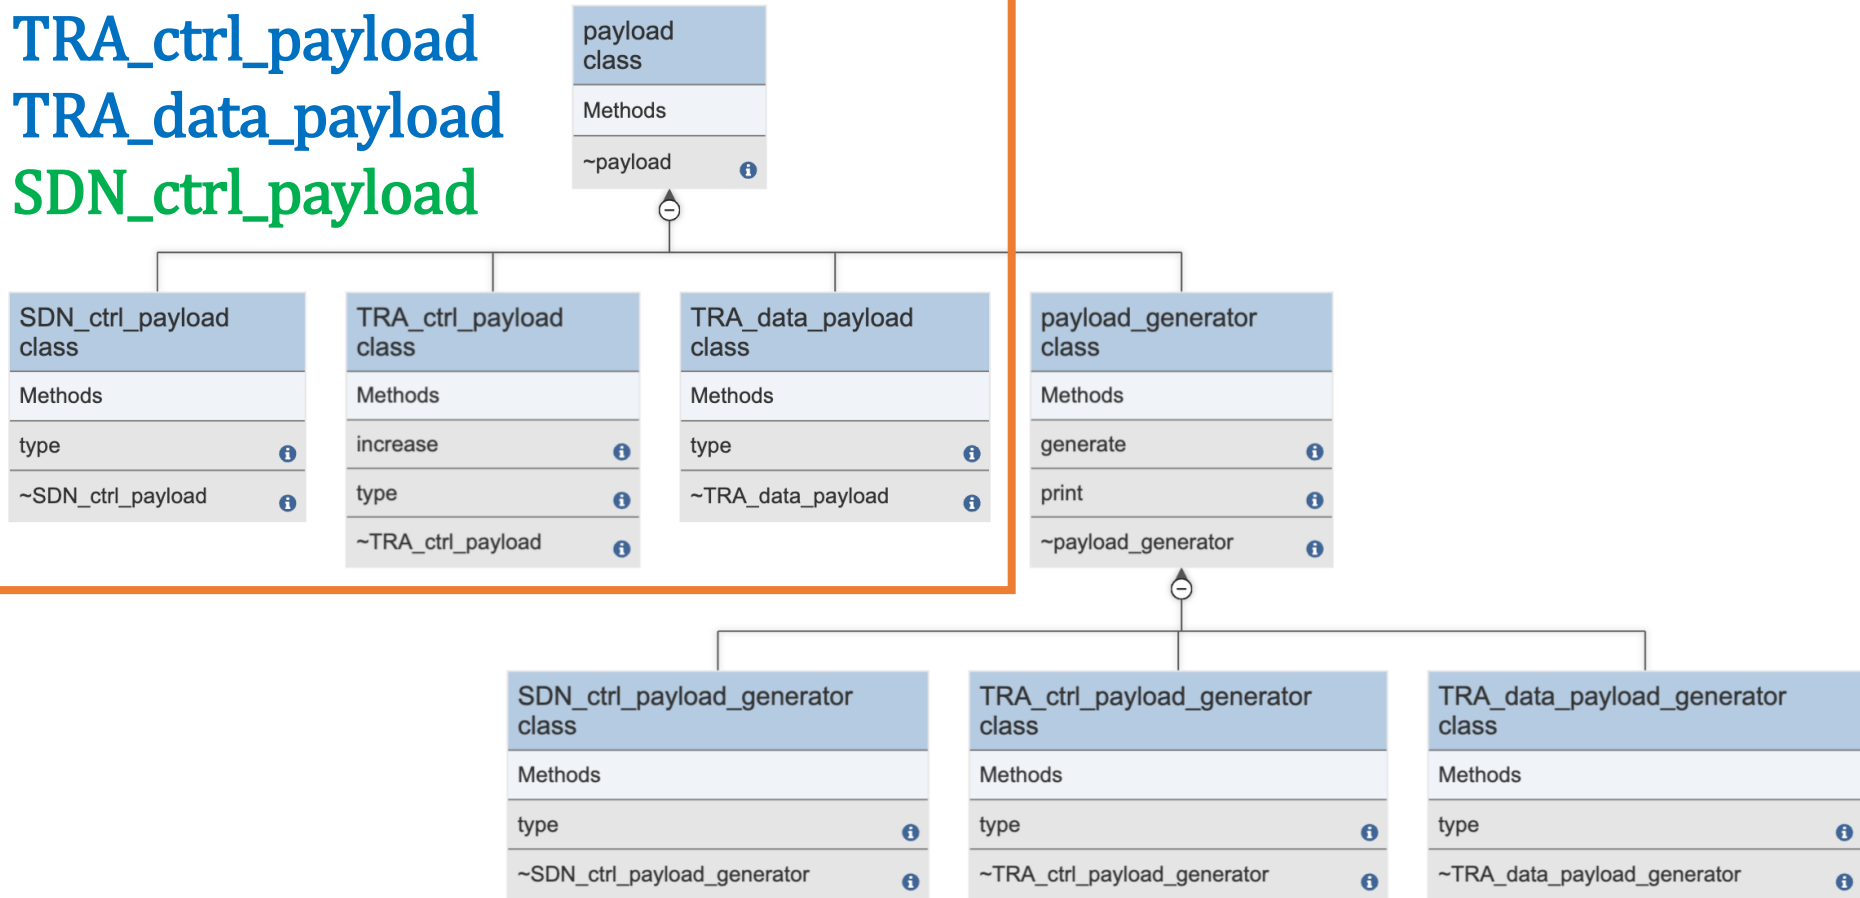
Inheritance

payload

TRA_ctrl_payload

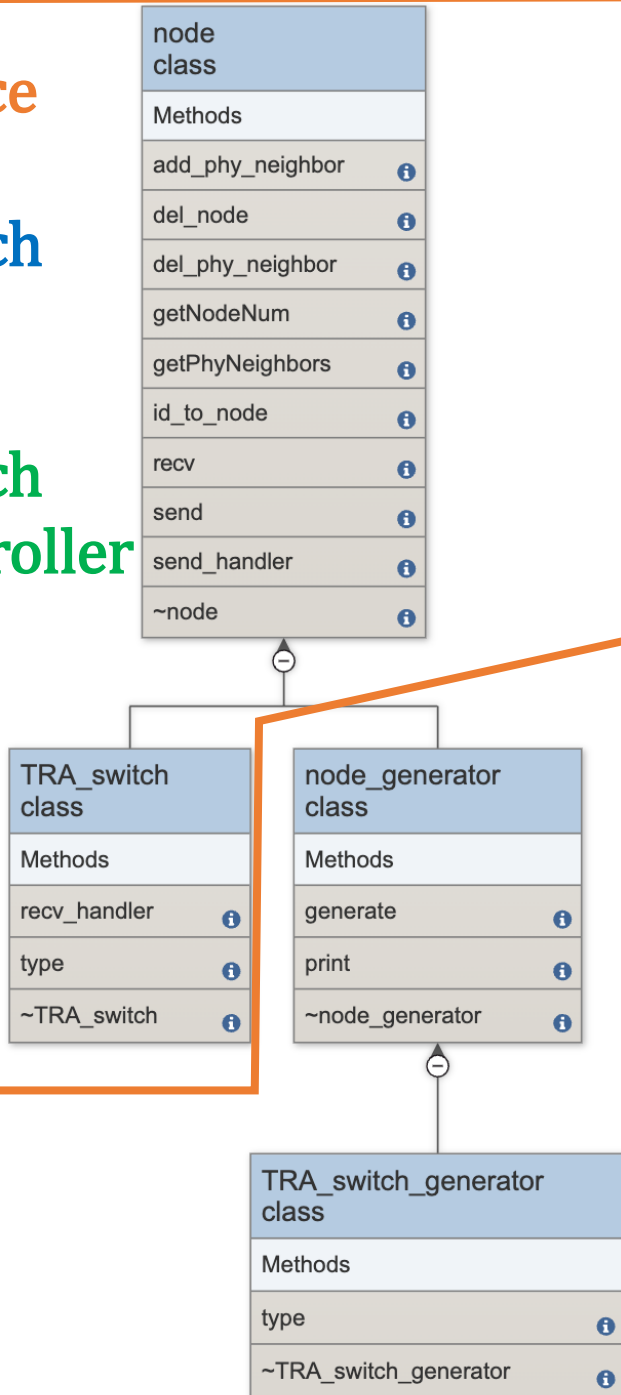
TRA_data_payload

SDN_ctrl_payload

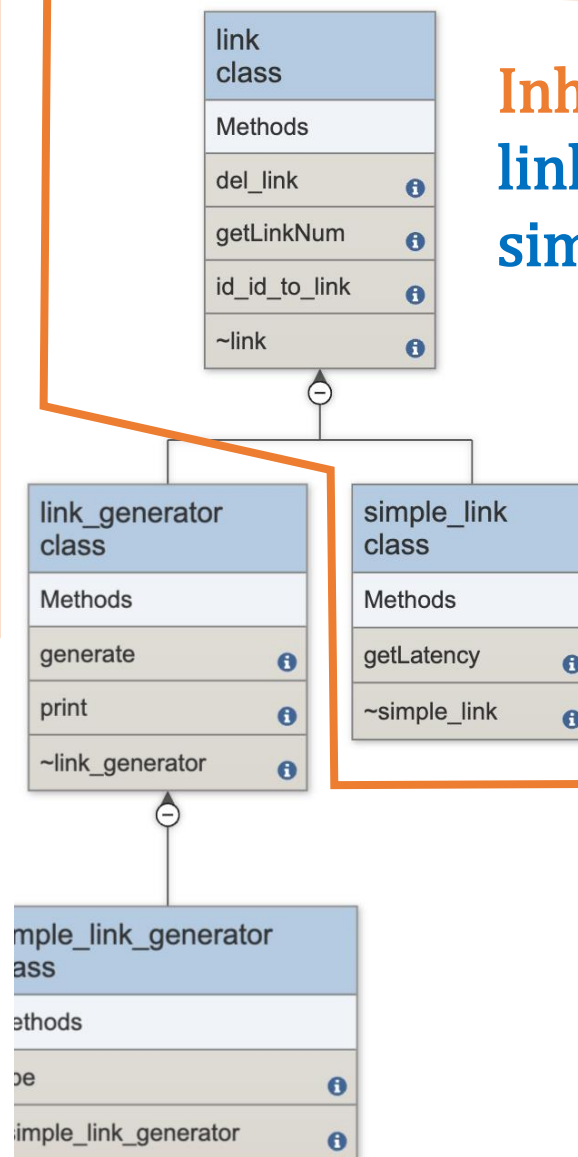


Inheritance
node
TRA_switch

Required:
SDN_switch
SDN_controller



Inheritance
link
simple_link



Input Sample:

use cin

Format:

#Nodes #Dsts #Links #Pairs SimTime SDN_CtrlTime Budget

DstID_List

NodeID UpgradeCost (BroadcastTime)


...

LinkID Node1 Node2

...

FlowID Src Dst FlowSize

...



15 1 28 3 300 200 120

0

0 40 100

1 30

2 40

3 50

4 40

5 60

6 40

7 40

8 20

9 50

10 50

11 20

12 30

13 20

14 30

0 0 5

1 0 10

2 0 12

3 0 14

4 1 2

5 1 3

6 1 4

7 2 3

8 2 10

9 2 11

10 3 4

11 3 6

12 3 7

13 4 8

14 4 9

15 5 6

16 5 7

17 5 9

18 5 10

19 5 14

20 6 7

21 6 10

22 7 9

23 8 9

24 9 13

25 10 12

26 11 12

27 13 14

0 1 0 3

1 2 0 4

2 3 0 5

Output Sample (not optimal):

use cout

You have to print the set of nodes that will be upgraded and the routing table for each node after event::start_simulate();

Then, you should implement recv_handler and some components

The remaining output will be automatically generated 😊

Note that the output could be different in different computers

use cout

Format:

UpgradedNodeIDList

DstID	NextID	Portion	NextID	Portion...
-------	--------	---------	--------	------------

[illegible]

2 9 14

0 0

0 2

0 10 50% 11 50%

0 2

0 9

0 0

0 5

0.5

0 0

0 12
12

13

14
2

0 0 70% 5 50%

Its own ID

Note

- Superb deadline: 5/16 Tue (from 4/18, you have 4 weeks)
- Deadline: 5/30 Tue (from 4/18, you have 6 weeks)
- Pass the test of our [online judge](#) platform
- Submit your code to [E-course2](#)
- Demonstrate your code [remotely](#) with TA
- **C++ Source code (only C++; compiled with g++)**
 - **Please use C++ library (i.e., no stdio, no stdlib)**
 - **Please use new and delete instead of malloc and free**
- Show a good programming style